

ejrzrvmcn

July 28, 2023

```
[20]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
[21]: df=pd.read_csv("/content/15_Horse Racing Results.csv")
df
```

```
[21]:
```

	Dato	Track	Race Number	Distance	Surface	Prize money	\
0	03.09.2017	Sha Tin	10	1400	Gress	1310000	
1	16.09.2017	Sha Tin	10	1400	Gress	1310000	
2	14.10.2017	Sha Tin	10	1400	Gress	1310000	
3	11.11.2017	Sha Tin	9	1600	Gress	1310000	
4	26.11.2017	Sha Tin	9	1600	Gress	1310000	
...	...	...	...	...	...	...	
27003	14.06.2020	Sha Tin	11	1200	Gress	1450000	
27004	21.06.2020	Sha Tin	2	1200	Gress	967000	
27005	21.06.2020	Sha Tin	4	1200	Gress	967000	
27006	21.06.2020	Sha Tin	5	1200	Gress	967000	
27007	21.06.2020	Sha Tin	11	1200	Gress	1450000	

	Starting position	Jockey	Jockey weight	Country	...	\
0	6	K C Leung	52	Sverige	...	
1	14	C Y Ho	52	Sverige	...	
2	8	C Y Ho	52	Sverige	...	
3	13	Brett Prebble	54	Sverige	...	
4	9	C Y Ho	52	Sverige	...	
...	...	...	...	...	...	
27003	6	A Hamelin	59	Australia	...	
27004	7	K C Leung	57	Australia	...	
27005	6	Blake Shinn	57	Australia	...	
27006	14	Joao Moreira	57	New Zealand	...	
27007	7	C Schofield	55	New Zealand	...	

	TrainerName	Race time	Path	Final place	FGrating	Odds	RaceType	\
0	CH Yip	83,38	2	9	110	22	Handicap	
1	CH Yip	81,56	3	4	124	48	Handicap	

2	CH Yip	82,36	1	6	118	11	Handicap
3	CH Yip	96,53	0	8	107	11	Handicap
4	CH Yip	94,17	0	3	123	40	Handicap
...	...	...	...	...	...	...	...
27003	WY So	70,87	1	9	104	25	Handicap
27004	KL Man	69,91	2	5	110	124	Handicap
27005	P O'Sullivan	69,49	0	3	114	88	Handicap
27006	AS Cruz	70,08	2	7	109	22	Handicap
27007	WY So	69,51	2	9	118	55	Handicap

	HorseId	JockeyId	TrainerID
0	1736	8656	6687
1	1736	8659	6687
2	1736	8659	6687
3	1736	8453	6687
4	1736	8659	6687
...	...	...	...
27003	29038	9111	6683
27004	29056	8656	6693
27005	29057	8778	6691
27006	29058	8443	6684
27007	29059	8655	6683

[27008 rows x 21 columns]

```
[22]: df.head()
```

```
[22]:
```

	Dato	Track	Race Number	Distance	Surface	Prize money	\
0	03.09.2017	Sha Tin	10	1400	Gress	1310000	
1	16.09.2017	Sha Tin	10	1400	Gress	1310000	
2	14.10.2017	Sha Tin	10	1400	Gress	1310000	
3	11.11.2017	Sha Tin	9	1600	Gress	1310000	
4	26.11.2017	Sha Tin	9	1600	Gress	1310000	

	Starting position	Jockey	Jockey weight	Country	...	TrainerName	\
0	6	K C Leung	52	Sverige	...	CH Yip	
1	14	C Y Ho	52	Sverige	...	CH Yip	
2	8	C Y Ho	52	Sverige	...	CH Yip	
3	13	Brett Prebble	54	Sverige	...	CH Yip	
4	9	C Y Ho	52	Sverige	...	CH Yip	

	Race time	Path	Final place	FGrating	Odds	RaceType	HorseId	JockeyId	\
0	83,38	2	9	110	22	Handicap	1736	8656	
1	81,56	3	4	124	48	Handicap	1736	8659	
2	82,36	1	6	118	11	Handicap	1736	8659	
3	96,53	0	8	107	11	Handicap	1736	8453	
4	94,17	0	3	123	40	Handicap	1736	8659	

```

      TrainerID
0         6687
1         6687
2         6687
3         6687
4         6687

```

```
[5 rows x 21 columns]
```

## 1 DATA CLEANING AND DATA PREPROCESSING

```
[23]: df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 27008 entries, 0 to 27007
Data columns (total 21 columns):
 #   Column                Non-Null Count  Dtype  
---  -
 0   Dato                  27008 non-null  object 
 1   Track                 27008 non-null  object 
 2   Race Number          27008 non-null  int64  
 3   Distance              27008 non-null  int64  
 4   Surface               27008 non-null  object 
 5   Prize money           27008 non-null  int64  
 6   Starting position     27008 non-null  int64  
 7   Jockey                27008 non-null  object 
 8   Jockey weight         27008 non-null  int64  
 9   Country               27008 non-null  object 
10  Horse age             27008 non-null  int64  
11  TrainerName           27008 non-null  object 
12  Race time             27008 non-null  object 
13  Path                  27008 non-null  int64  
14  Final place           27008 non-null  int64  
15  FGGrating             27008 non-null  int64  
16  Odds                  27008 non-null  object 
17  RaceType              27008 non-null  object 
18  HorseId               27008 non-null  int64  
19  JockeyId              27008 non-null  int64  
20  TrainerID             27008 non-null  int64  
dtypes: int64(12), object(9)
memory usage: 4.3+ MB

```

```
[24]: df.describe()
```

```
[24]:
```

	Race Number	Distance	Prize money	Starting position	\
count	27008.000000	27008.000000	2.700800e+04	27008.000000	
mean	5.268624	1401.666173	1.479445e+06	6.741447	
std	2.780088	276.065045	2.162109e+06	3.691071	
min	1.000000	1000.000000	6.600000e+05	1.000000	
25%	3.000000	1200.000000	9.200000e+05	4.000000	
50%	5.000000	1400.000000	9.670000e+05	7.000000	
75%	8.000000	1650.000000	1.450000e+06	10.000000	
max	11.000000	2400.000000	2.800000e+07	14.000000	

	Jockey weight	Horse age	Path	Final place	FGrating	\
count	27008.000000	27008.000000	27008.000000	27008.000000	27008.000000	
mean	55.867373	5.246408	1.678021	6.685834	113.428318	
std	2.737006	1.519880	1.631784	3.664551	13.314949	
min	47.000000	2.000000	0.000000	1.000000	-5.000000	
25%	54.000000	4.000000	0.000000	4.000000	106.000000	
50%	56.000000	5.000000	1.000000	7.000000	114.000000	
75%	58.000000	6.000000	3.000000	10.000000	122.000000	
max	63.000000	12.000000	11.000000	14.000000	157.000000	

	HorseId	JockeyId	TrainerID
count	27008.000000	27008.000000	27008.000000
mean	23904.874889	8586.732042	6668.559205
std	2028.860311	569.616932	79.978067
min	1736.000000	227.000000	6431.000000
25%	22364.000000	8651.000000	6683.000000
50%	22868.500000	8658.000000	6687.000000
75%	25310.000000	8663.000000	6693.000000
max	29059.000000	9113.000000	7004.000000

```
[25]: df.columns
```

```
[25]: Index(['Dato', 'Track', 'Race Number', 'Distance', 'Surface', 'Prize money',
          'Starting position', 'Jockey', 'Jockey weight', 'Country', 'Horse age',
          'TrainerName', 'Race time', 'Path', 'Final place', 'FGrating', 'Odds',
          'RaceType', 'HorseId', 'JockeyId', 'TrainerID'],
          dtype='object')
```

```
[26]: df1=df.dropna(axis=1)
df1
```

```
[26]:
```

	Dato	Track	Race Number	Distance	Surface	Prize money	\
0	03.09.2017	Sha Tin	10	1400	Gress	1310000	
1	16.09.2017	Sha Tin	10	1400	Gress	1310000	
2	14.10.2017	Sha Tin	10	1400	Gress	1310000	
3	11.11.2017	Sha Tin	9	1600	Gress	1310000	
4	26.11.2017	Sha Tin	9	1600	Gress	1310000	

...	...	...	...	...	...	...	...
27003	14.06.2020	Sha Tin	11	1200	Gress	1450000	
27004	21.06.2020	Sha Tin	2	1200	Gress	967000	
27005	21.06.2020	Sha Tin	4	1200	Gress	967000	
27006	21.06.2020	Sha Tin	5	1200	Gress	967000	
27007	21.06.2020	Sha Tin	11	1200	Gress	1450000	

	Starting position	Jockey	Jockey weight	Country	...	\
0	6	K C Leung	52	Sverige	...	
1	14	C Y Ho	52	Sverige	...	
2	8	C Y Ho	52	Sverige	...	
3	13	Brett Prebble	54	Sverige	...	
4	9	C Y Ho	52	Sverige	...	

...	...	...	...	...	...	...
27003	6	A Hamelin	59	Australia	...	
27004	7	K C Leung	57	Australia	...	
27005	6	Blake Shinn	57	Australia	...	
27006	14	Joao Moreira	57	New Zealand	...	
27007	7	C Schofield	55	New Zealand	...	

	TrainerName	Race time	Path	Final place	FGrating	Odds	RaceType	\
0	CH Yip	83,38	2	9	110	22	Handicap	
1	CH Yip	81,56	3	4	124	48	Handicap	
2	CH Yip	82,36	1	6	118	11	Handicap	
3	CH Yip	96,53	0	8	107	11	Handicap	
4	CH Yip	94,17	0	3	123	40	Handicap	
...	...	...	...	...	...	...	...	
27003	WY So	70,87	1	9	104	25	Handicap	
27004	KL Man	69,91	2	5	110	124	Handicap	
27005	P O'Sullivan	69,49	0	3	114	88	Handicap	
27006	AS Cruz	70,08	2	7	109	22	Handicap	
27007	WY So	69,51	2	9	118	55	Handicap	

	HorseId	JockeyId	TrainerID
0	1736	8656	6687
1	1736	8659	6687
2	1736	8659	6687
3	1736	8453	6687
4	1736	8659	6687
...	...	...	...
27003	29038	9111	6683
27004	29056	8656	6693
27005	29057	8778	6691
27006	29058	8443	6684
27007	29059	8655	6683

[27008 rows x 21 columns]

```
[27]: df1.columns
```

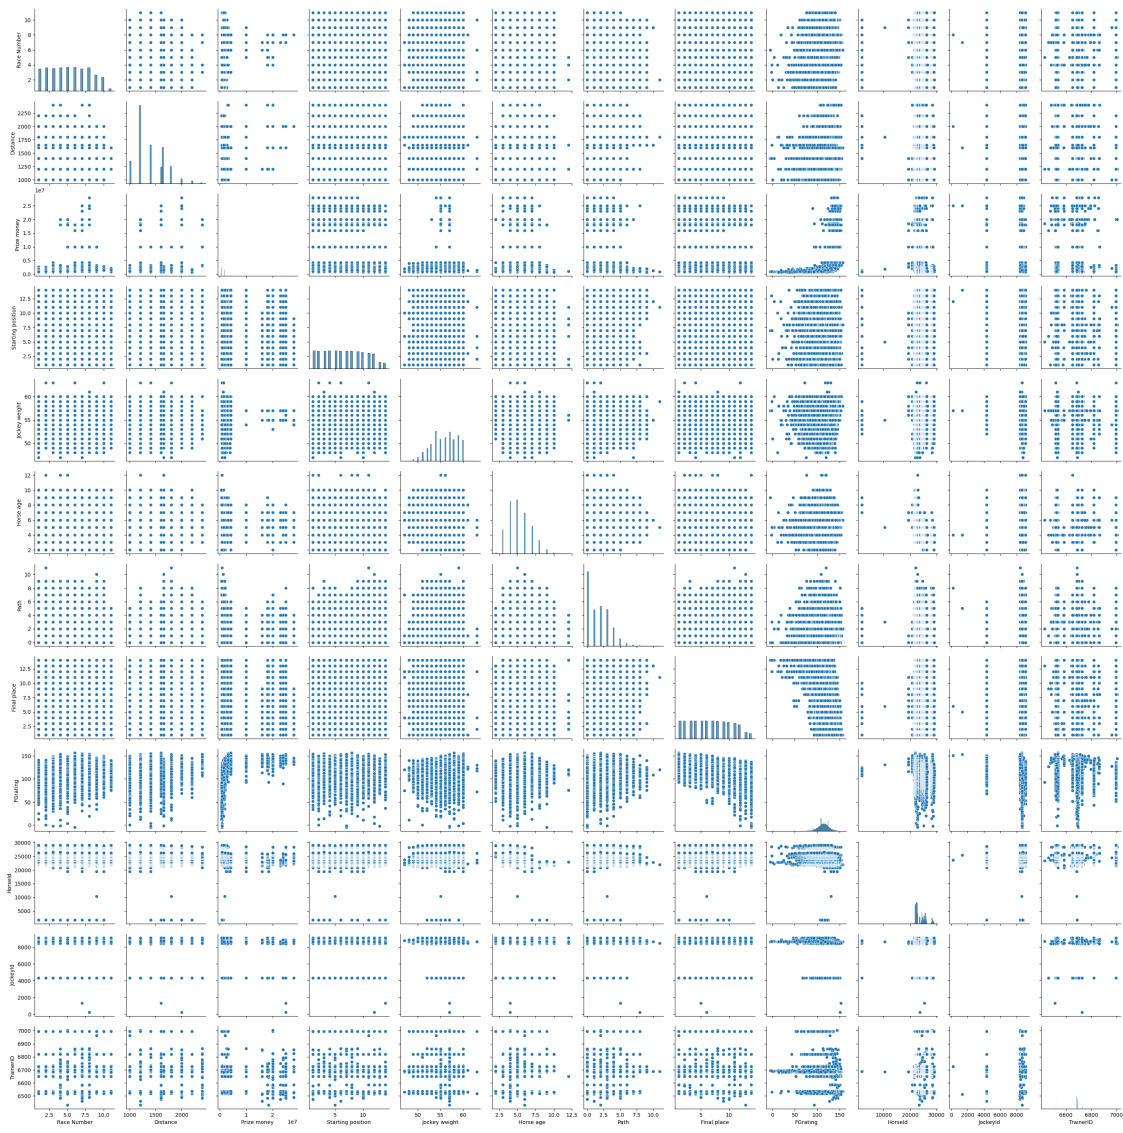
```
[27]: Index(['Dato', 'Track', 'Race Number', 'Distance', 'Surface', 'Prize money',  
         'Starting position', 'Jockey', 'Jockey weight', 'Country', 'Horse age',  
         'TrainerName', 'Race time', 'Path', 'Final place', 'FGrating', 'Odds',  
         'RaceType', 'HorseId', 'JockeyId', 'TrainerID'],  
        dtype='object')
```

```
[28]: df1=df1[['Dato', 'Track', 'Race Number', 'Distance', 'Surface', 'Prize money',  
             'Starting position', 'Jockey', 'Jockey weight', 'Country', 'Horse age',  
             'TrainerName', 'Race time', 'Path', 'Final place', 'FGrating', 'Odds',  
             'RaceType', 'HorseId', 'JockeyId', 'TrainerID']]
```

## 2 EDA AND VISUALIZATION

```
[29]: sns.pairplot(df1)
```

```
[29]: <seaborn.axisgrid.PairGrid at 0x7ea8e39bbb80>
```



```
[30]: sns.distplot(df1['Distance'])
```

<ipython-input-30-55baf5480dab>:1: UserWarning:


`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(df1['Distance'])
```

```
[30]: <Axes: xlabel='Distance', ylabel='Density'>
```




```
[31]: sns.heatmap(df1.corr())
```

```
<ipython-input-31-3ed1a1a51dc0>:1: FutureWarning: The default value of  
numeric_only in DataFrame.corr is deprecated. In a future version, it will  
default to False. Select only valid columns or specify the value of numeric_only  
to silence this warning.  
sns.heatmap(df1.corr())
```

```
[31]: <Axes: >
```





### 3 TO TRAIN THE MODEL AND MODEL BUILDING

```
[32]: x=df[['Race Number', 'Prize money',
           'Starting position', 'Jockey weight', 'Horse age', 'Final place',
           'FGating', 'Horseld', 'JockeyId', 'TrainerID']]
      y=df['Distance']
```

```
[33]: from sklearn.model_selection import train_test_split
      x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

```
[34]: from sklearn.linear_model import LinearRegression
      lr=LinearRegression()
      lr.fit(x_train,y_train)
```

```
[34]: LinearRegression()
```

```
[35]: lr.intercept_
```

[35]: 1702.2246084652907


```
[36]: coeff=pd.DataFrame(lr.coef_,x.columns,columns=['Co-efficient'])
      coeff
```

```
[36]:
```

	Co-efficient
Race Number	-6.066791
Prize money	0.000015
Starting position	-0.576600
Jockey weight	-1.939342
Horse age	22.775448
Final place	6.388438
FGrating	3.067410
HorseId	-0.006693
JockeyId	-0.006062
TrainerID	-0.071689

```
[37]: prediction =lr.predict(x_test)
      plt.scatter(y_test,prediction)
```

[37]: <matplotlib.collections.PathCollection at 0x7ea8d2eb3ee0>



## 4 ACCURACY

```
[38]: lr.score(x_test,y_test)
```

```
[38]: 0.06099752116422741
```

```
[39]: lr.score(x_train,y_train)
```

```
[39]: 0.06411986650423529
```

```
[40]: from sklearn.linear_model import Ridge,Lasso
```

```
[41]: rr=Ridge(alpha=10)  
      rr.fit(x_train,y_train)
```

```
[41]: Ridge(alpha=10)
```

```
[42]: rr.score(x_test,y_test)
```

```
[42]: 0.06099829608026752
```

```
[43]: rr.score(x_train,y_train)
```

```
[43]: 0.06411986523298252
```

```
[44]: la=Lasso(alpha=10)  
      la.fit(x_train,y_train)
```

```
[44]: Lasso(alpha=10)
```

```
[45]: la.score(x_test,y_test)
```

```
[45]: 0.05982733108440974
```

```
[46]: la.score(x_train,y_train)
```

```
[46]: 0.06224170620823566
```